

Initial Decision Matrix: Technical Architecture

This document tracks the core technical choices made at the start of the project, including the tradeoffs considered and the final path selected.

Feature	Option A	Option B	Selected Path	Engineering Reasoning
Orchestration	LangChain / Frameworks	Scratch-built SDK implementation	Scratch (Gemini SDK)	Shows deeper understanding of the message loop and state management. Better for a "Quant Engineer" role to demonstrate logic control.
SQL Access	Text-to-SQL (Open)	Pre-defined Functions (Rigid)	Validated Text-to-SQL	A hybrid "middle ground." Allows flexibility for "unseen queries" but uses a Python-based validator to block destructive commands and enforce limits.
Vector DB	Managed Cloud (Pinecone)	Local-first (ChromaDB)	ChromaDB	Industry standard for local/containerized RAG. Avoids 3rd party cloud dependencies while demonstrating high-quality retrieval architecture.
Access Model	Full Context Exposure	Tiered Access Masking	Tiered Masking	The LLM is treated as a "masked" entity with tool-only access. The User and Backend have full visibility, ensuring the model can't leak the entire dataset.
Database	PostgreSQL	SQLite	SQLite	Priority on "Deployment Readiness" and ease of evaluation. One-file DB means the interviewer can run the app instantly without DB config friction.

Post-Decision Requirements

1. **Trace View:** Must map LLM outputs back to specific PDF page/snippet IDs without exposing the whole doc to the model in the prompt.

2. **Double-Request Guard:** Implement debouncing/loading states on the UI to prevent "chaotic" user input issues.
3. **Data Consistency:** Mock data generation script must ensure the CSVs and PDFs tell the same story (e.g., matching titles and trends).