

NeonPlay Media
internal policy + ai usage notes

Prepared by: ops / infra / whoever had latest version
Internal only pls
version 0.9.7 draft

purpose

This doc is supposed to outline:

- internal data handling
- AI tool restrictions
- analytics access
- document usage rules
- basic security expectations

Some sections still being finalized with legal/compliance.

general rules

Employees should:

- avoid sharing internal exports publicly
- not upload customer data into random AI tools
- report suspicious access/activity
- use company-approved dashboards where possible

Current governance maturity still kinda early so a lot of workflows are being standardized rn.

data access

low risk data

stuff like:

- campaign summaries
- aggregated metrics
- public marketing assets

usually accessible internally.

restricted data

includes:

- customer emails
- payment info
- auth/session logs
- raw subscriber exports
- internal finance docs

This data should stay restricted to approved people only.

AI tools definitely should NOT get raw subscriber exports.

analytics access

Current setup:

- growth team has dashboard access
- product has engagement metrics
- engineering has infra metrics
- finance exports limited to finance team

Some access overlap still exists because permissions were initially set up manually.

cleanup in progress.

AI tooling

Internal AI tools currently allowed for:

- summarizing reports
- helping with analytics retrieval
- trend explanations
- internal search

Not allowed:

- direct database writes
- raw customer exposure
- payment data visibility
- bypassing permissions

Most AI queries currently routed through retrieval APIs instead of direct db access.

logging

We currently log:

- query timestamps
- user access
- dashboard exports
- retrieval activity
- failed auth attempts

Need better centralized logging though because some systems still logging separately.

password/security notes

Please stop sharing staging passwords in Slack.

yes this had to be added because it happened twice.

document handling

Supported uploads rn:

- pdf
- csv
- xlsx
- txt

DOCX support exists but parsing still inconsistent in some cases.

Corrupted files should:

- fail safely
- not crash ingestion
- get isolated for review

There was an issue last month where one malformed csv broke a nightly analytics job for ~40 mins.

retention rules

Current retention windows:

Data Type	Retention
watch logs	~18 months
analytics exports	12 months
audit logs	longer probably
support chats	TBD

Need legal signoff on final retention policy.

audit logging

Systems should ideally track:

- who accessed what
- when exports happened
- what files were uploaded
- retrieval queries
- failed access attempts

Some older services still missing full audit coverage.

access issues

Known problems:

- duplicated permissions
- old contractor accounts not removed fast enough
- inconsistent dashboard visibility
- staging/prod confusion in some internal tools

infra team already working on cleanup.

external tools

Before uploading internal files to external services:

- check with ops/security
- avoid customer exports
- avoid finance docs
- verify data masking if possible

One team uploaded a raw csv to a public notebook tool earlier this year so now everyone gets reminders about this constantly.

AI safety stuff

Current restrictions:

- AI cannot directly expose raw customer tables
- restricted documents should stay masked
- internal retrieval logs should remain traceable

Need clearer policy wording around embeddings/vector storage later.

still unfinished.

incident handling

If any of these happen:

- exposed customer info
- suspicious exports
- unexpected AI outputs
- broken access controls
- leaked docs

report it ASAP to:

- infra
- ops
- whoever is on-call honestly

Current incident process still kinda informal.

infra/security notes

Things infra team monitoring:

- unusual export spikes
- repeated failed logins
- large retrieval requests
- abnormal API usage
- random scraping attempts

One internal dashboard got hammered accidentally after a bad retry loop during May.

governance gaps

Things still not fully standardized:

- naming conventions
- campaign metadata
- dashboard ownership
- AI usage approvals
- retention enforcement

Company grew faster than internal governance processes.

temporary rules added recently

AI usage

Do NOT:

- paste raw subscriber exports into public AI chat tools
- share internal reports externally
- expose finance sheets in demos

exports

Please label exports properly.

We currently have files named:

- final_v2_REAL.csv
- final_final2.csv
- q2numbers_latest_USETHIS.csv

which is becoming a problem.

roadmap for cleanup

Planned improvements:

1. centralized permissions
 2. proper audit coverage
 3. cleaner retention policies
 4. unified dashboard naming
 5. AI governance review process
-

final notes

This doc is still evolving as internal systems mature.

Main goals:

- avoid accidental exposure
- improve governance
- support fast-moving teams without blocking work
- make AI tooling safer + more traceable

Prepared by: ops + infra + analytics

NeonPlay Media

END

note:

*sections merged from multiple docs after old policy file export
broke during migration from google drive*